

Lecture 5

Jan 11, 2024

fork-exec model:

`fork()` creates a copy of the current process

`exec*()` replaces the current process' code and address space with the code for a different program

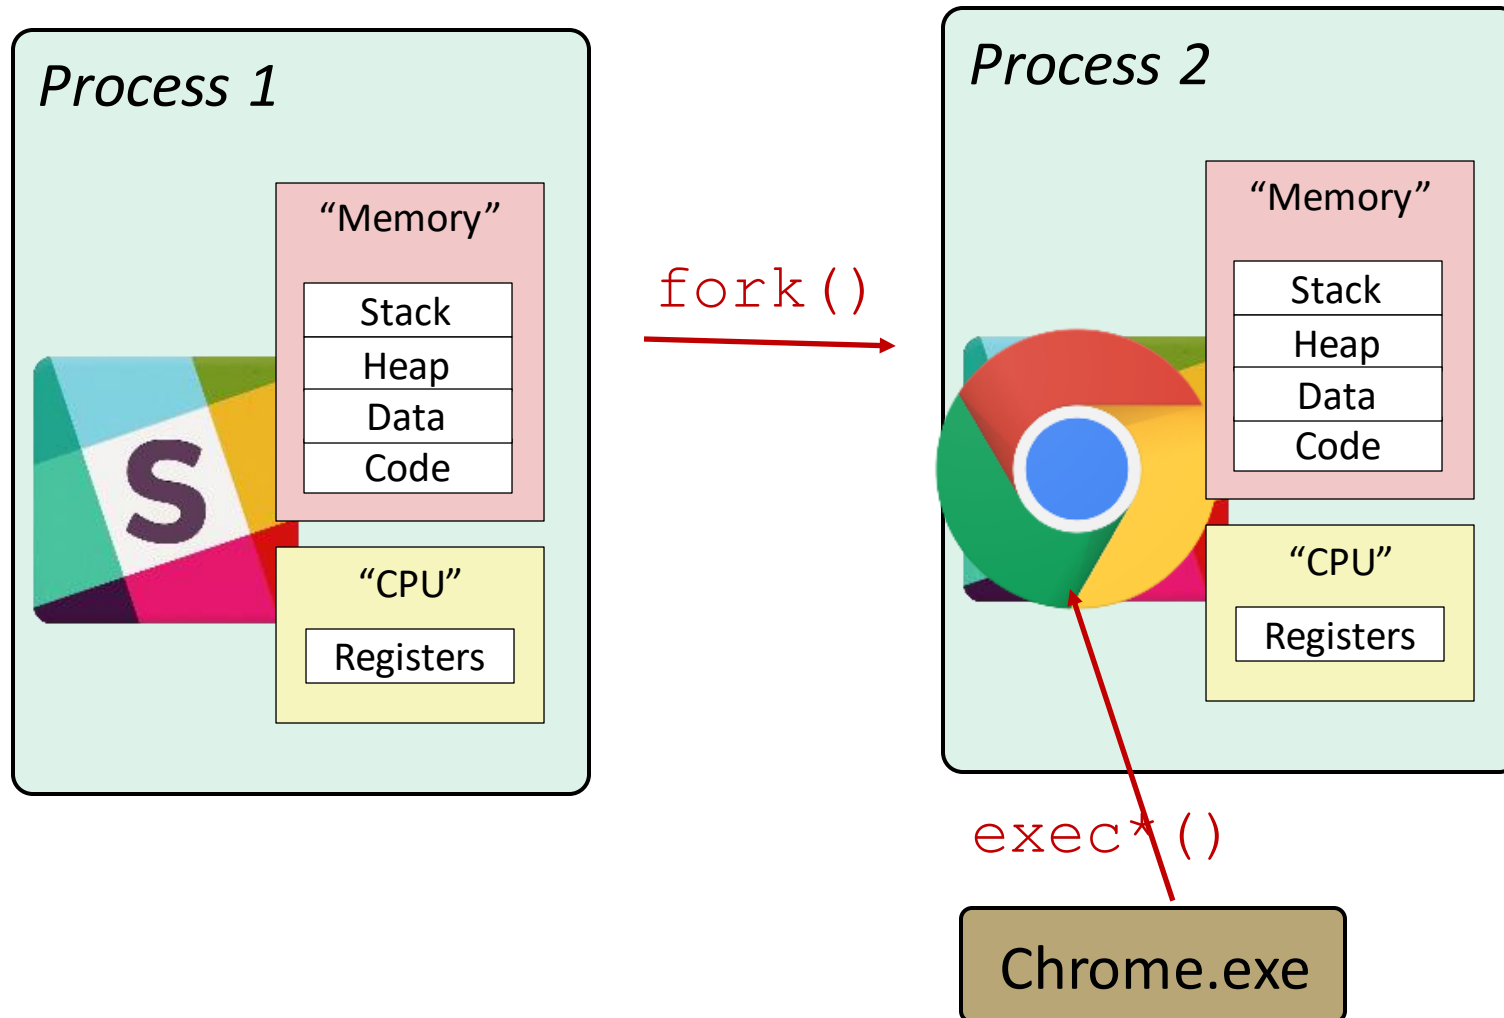
- Whole family of `exec` calls – see **`exec(3)`** and **`execve(2)`**

```
// Example arguments: path="/usr/bin/ls",  
//      argv[0]="/usr/bin/ls", argv[1]="-ahl", argv[2]=NULL  
void fork_exec(char *path, char *argv[]) {  
    pid_t pid = fork();  
    if (pid != 0) {  
        printf("Parent: created a child %d\n", pid);  
    } else {  
        printf("Child: about to exec a new program\n");  
        execv(path, argv);  
    }  
    printf("This line printed by parent only!\n");  
}
```

What happens during exec?

- •After fork, parent and child are running same code
- –Not too useful!
- •A process can run `exec()` to load another executable to its memory image
- –So, a child can run a different program from parent
- •Variants of `exec()`, e.g., to pass command line arguments to new executable
- On Linux, there are six variants of `exec()`: `execl`, `execvp()`, `execle()`, `execv()`, `execvp()`, and `execvpe()`.
- Read the man pages to learn more.

Exec()

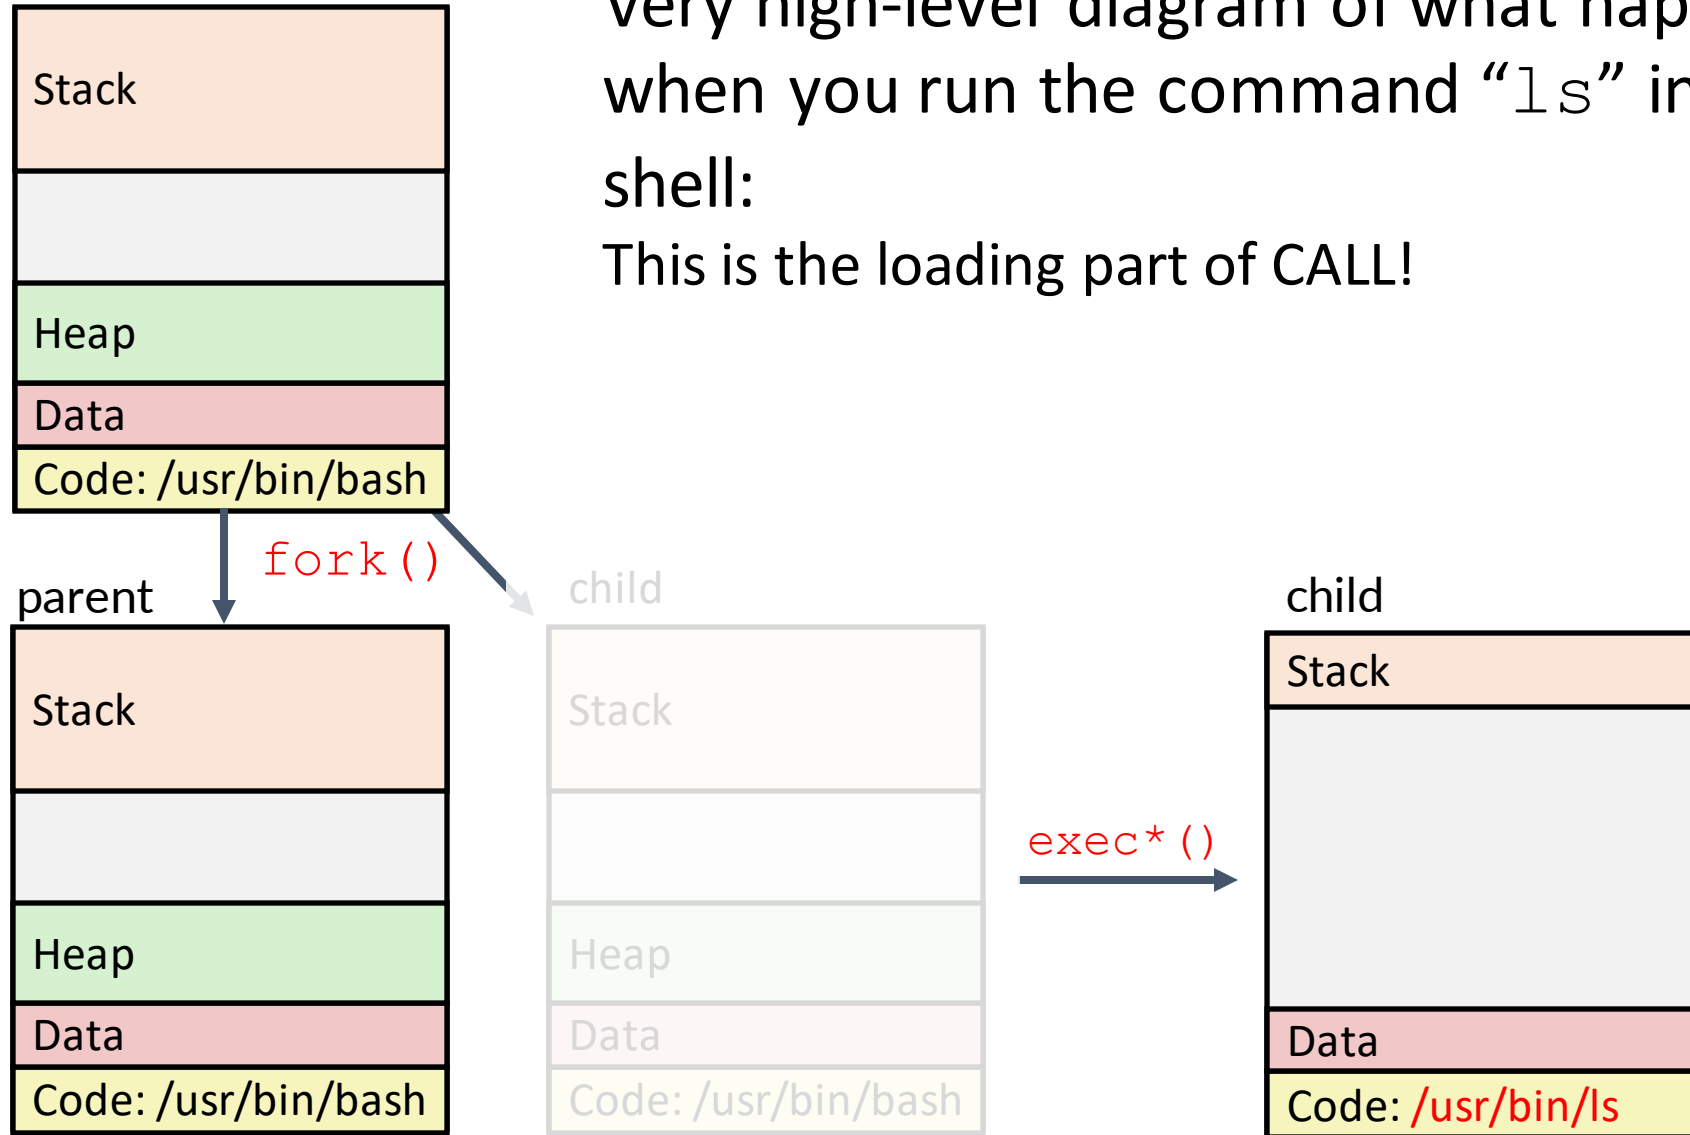


Why do we need fork() and exec()

- In a basic OS, the init process is created after initialization of hardware
- The init process spawns a shell like bash
- Shell reads user command, forks a child, execs the command executable, waits for it to finish, and reads next command
- Common commands like ls are all executables that are simply exec'ed by the shell

Very high-level diagram of what happens when you run the command “ls” in a Linux shell:

This is the loading part of CALL!



```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <sys/wait.h>
6
7  int main(int argc, char *argv[]) {
8      printf("hello world (pid:%d)\n", (int) getpid());
9      int rc = fork();
10     if (rc < 0) {                // fork failed; exit
11         fprintf(stderr, "fork failed\n");
12         exit(1);
13     } else if (rc == 0) { // child (new process)
14         printf("hello, I am child (pid:%d)\n", (int) getpid());
15         char *myargs[3];
16         myargs[0] = strdup("wc"); // program: "wc" (word count)
17         myargs[1] = strdup("p3.c"); // argument: file to count
18         myargs[2] = NULL;          // marks end of array
19         execvp(myargs[0], myargs); // runs word count
20         printf("this shouldn't print out");
21     } else {                      // parent goes down this path (main)
22         int rc_wait = wait(NULL);
23         printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n",
24               rc, rc_wait, (int) getpid());
25     }
26     return 0;
27 }

```

prompt> ./p3

hello world (pid:29383)

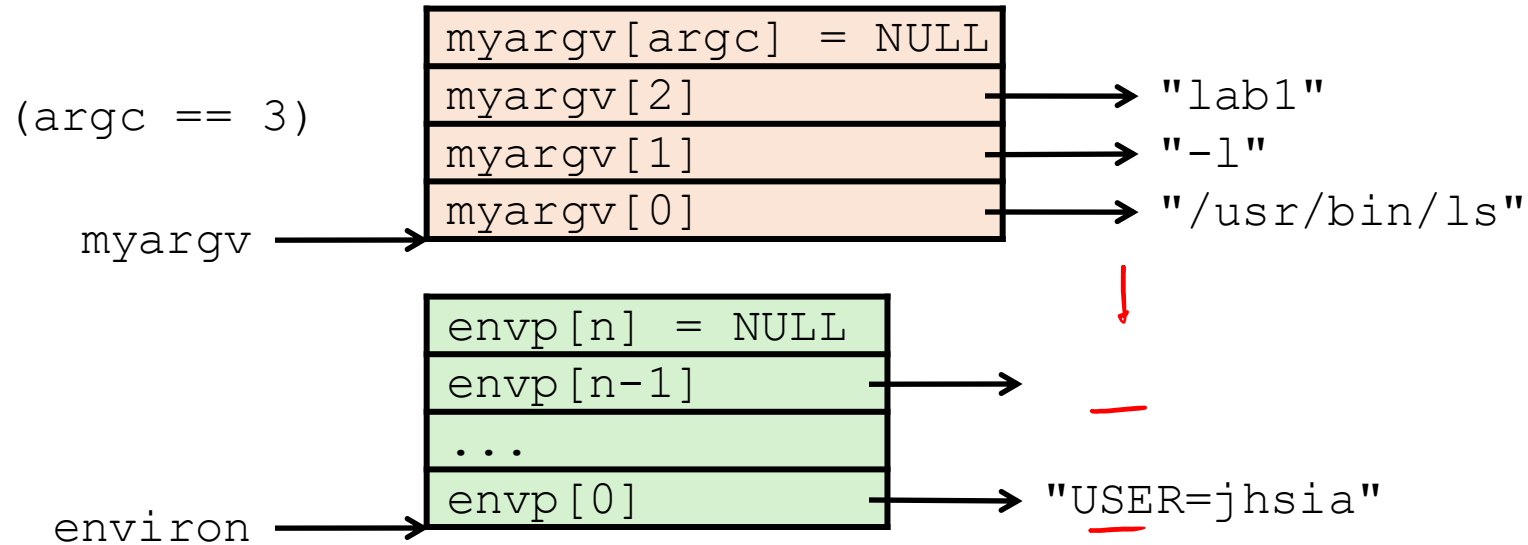
hello, I am child (pid:29384)

29 107 1030 p3.c

hello, I am parent of 29384 (rc_wait:29384) (pid:29383)

prompt>

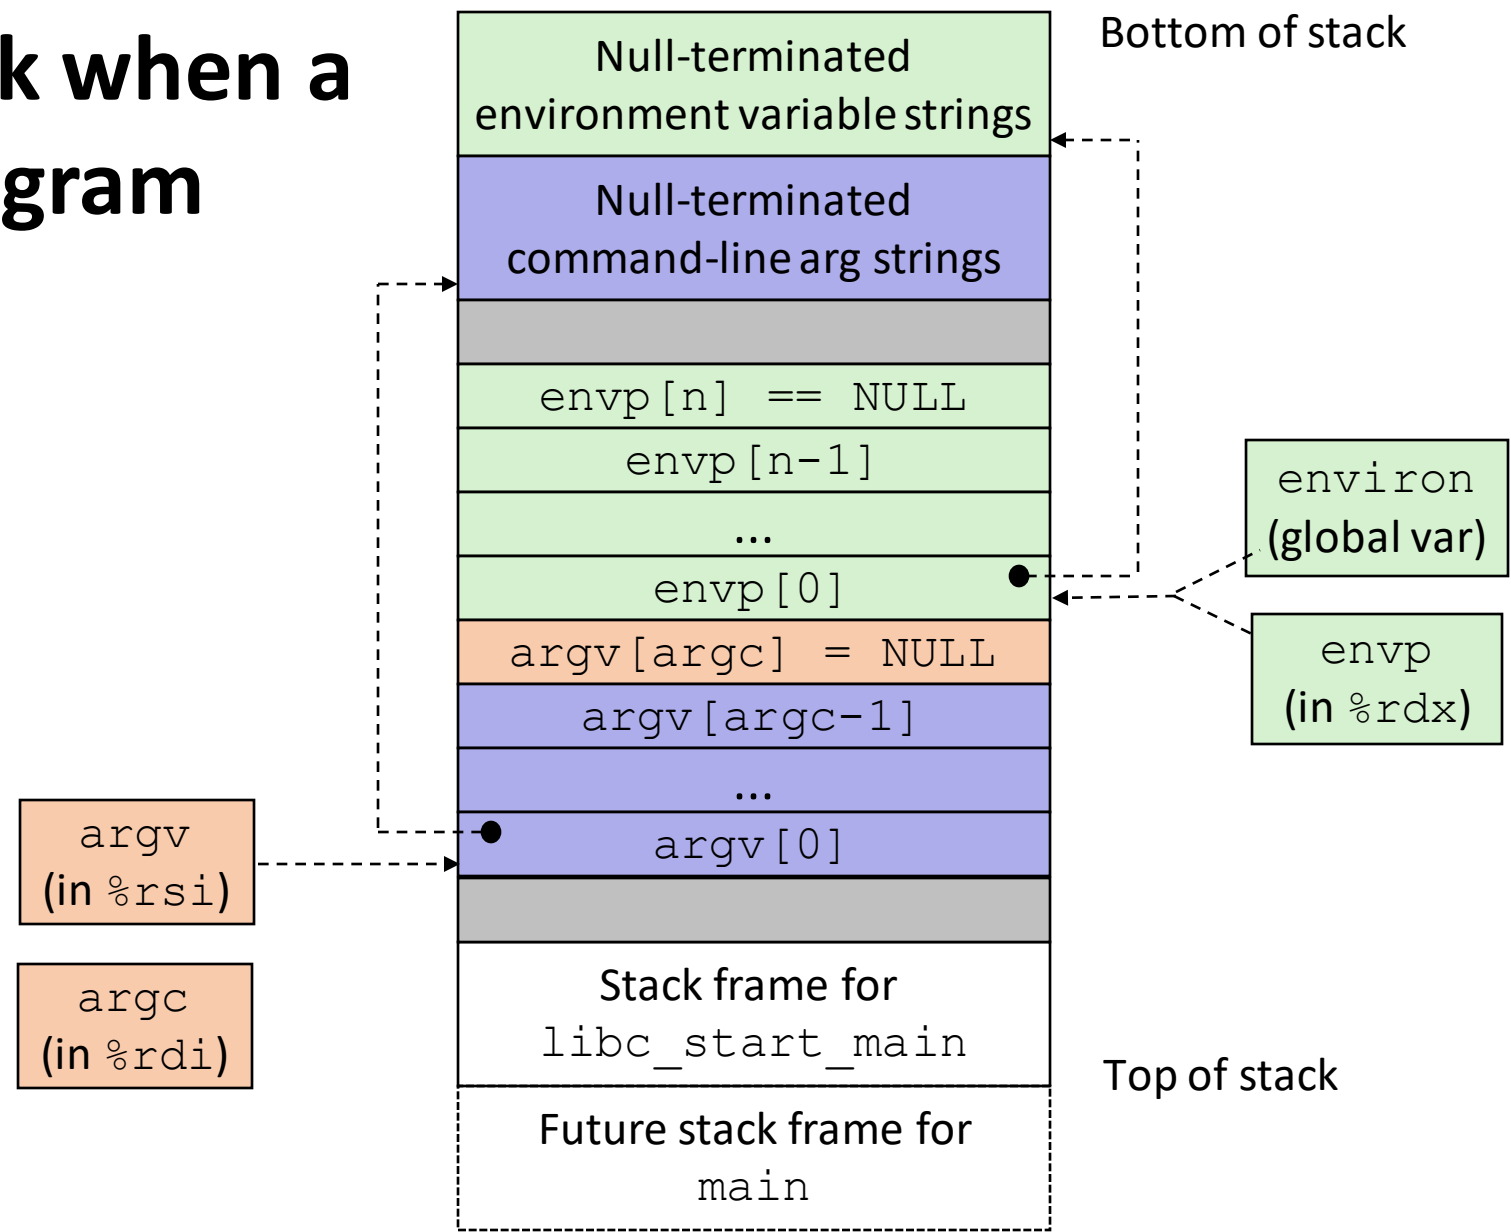
Execute `"/usr/bin/ls -l lab1"` in child process using current environment:



```
if ((pid = fork()) == 0) { /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```

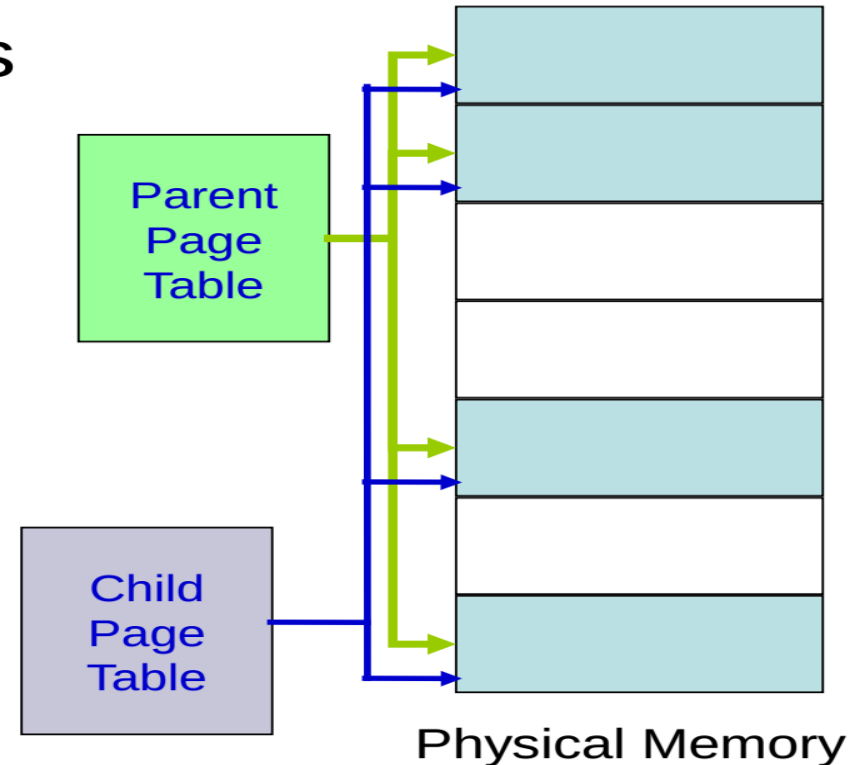
Run the `printenv` command in a Linux shell to see your own environment variables

Structure of the Stack when a new program starts



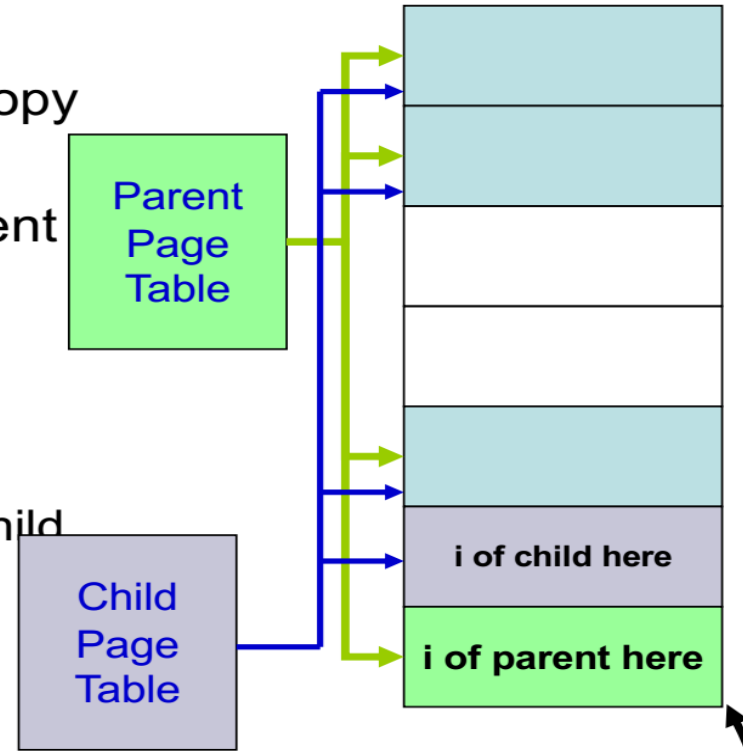
Virtual Addressing Advantage (easy to make copies of a process)

- Making a copy of a process is called forking.
 - Parent (is the original)
 - child (is the new process)
- When fork is invoked,
 - child is an exact copy of parent
 - When fork is called all pages are shared between parent and child
 - Easily done by copying the parent's page tables



Copy on Write (CoW)

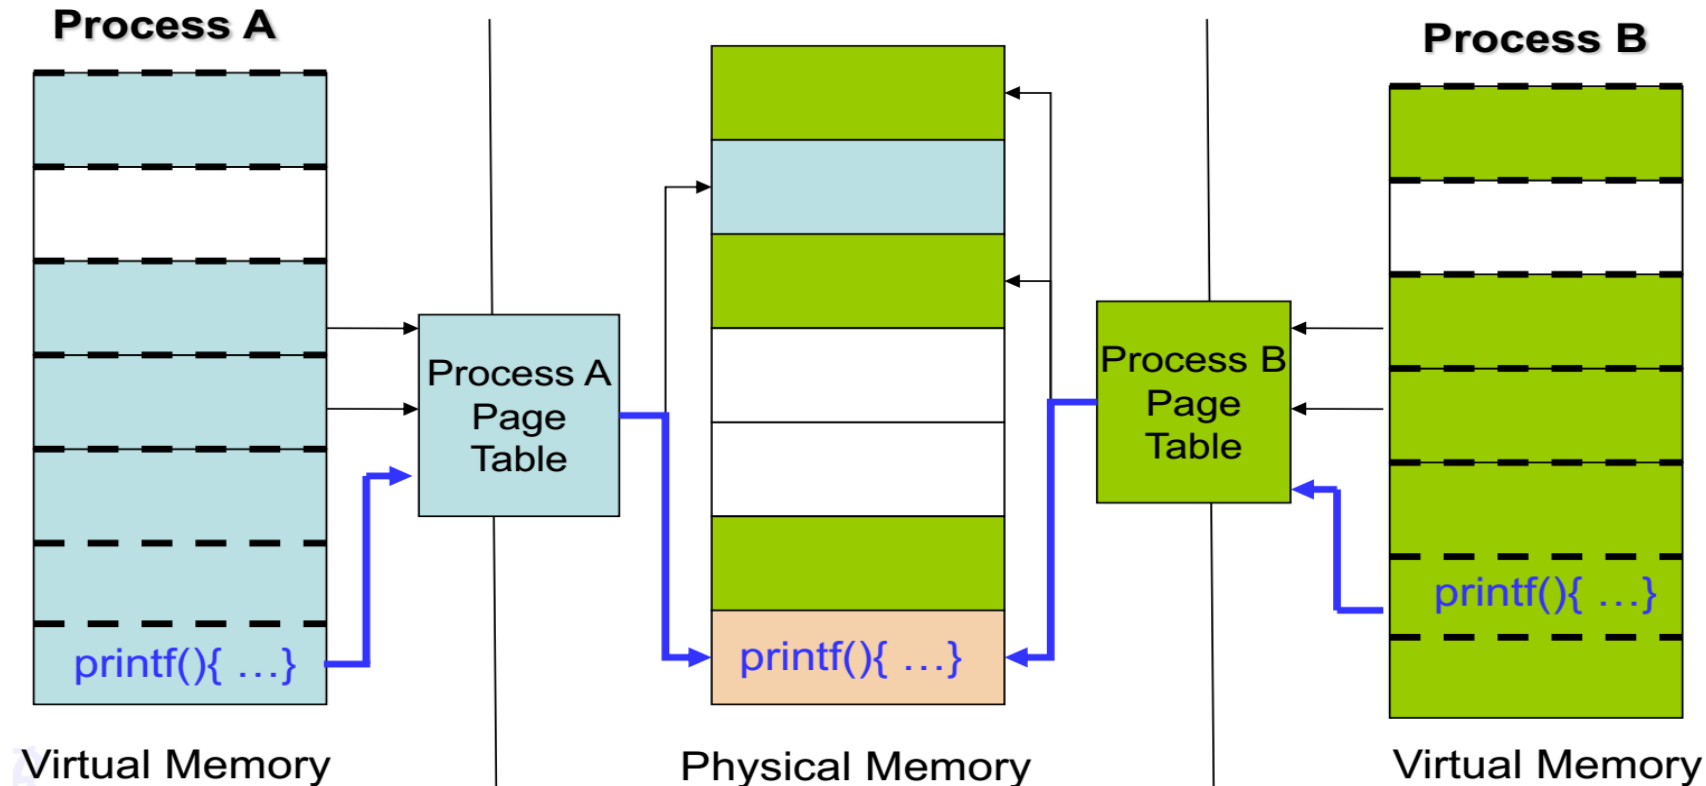
- When data in any of the shared pages change, OS intercepts and makes a copy of the page.
- Thus, parent and child will have different copies of **this** page
- **Why?**
 - A large portion of executables are not used.
 - Copying each page from parent and child would incur significant disk swapping.. huge performance penalties.
 - Postpone coping of pages as much as possible thus optimizing performance



This page now is no longer shared

Virtual Addressing Advantages (Shared libraries)

- Many common functions such as *printf* implemented in shared libraries
- Pages from shared libraries, shared between processes



How COW works??

- When forking,
 - – Kernel makes COW pages as read only
 - – Any write to the pages would cause a page fault
 - – The kernel detects that it is a COW page and duplicates the page

More details on Shell

- Shell can manipulate the child in strange ways.
- Suppose you want to redirect output from a command to a file
- **prompt>ls > foo.txt**
- Shell spawns a child, rewires its standard output to a file, then calls exec on the child

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <fcntl.h>
6  #include <sys/wait.h>
7
8  int main(int argc, char *argv[]) {
9      int rc = fork();
10     if (rc < 0) {
11         // fork failed
12         fprintf(stderr, "fork failed\n");
13         exit(1);
14     } else if (rc == 0) {
15         // child: redirect standard output to a file
16         close(STDOUT_FILENO);
17         open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU);
18
19         // now exec "wc"...
20         char *myargs[3];
21         myargs[0] = strdup("wc"); // program: wc (word count)
22         myargs[1] = strdup("p4.c"); // arg: file to count
23         myargs[2] = NULL; // mark end of array
24         execvp(myargs[0], myargs); // runs word count
25     } else {
26         // parent goes down this path (main)
27         int rc_wait = wait(NULL);
28     }
29     return 0;
30 }

```

prompt> ./p4

prompt> cat p4.output

32 109 846 p4.c

prompt>

So, should we rewrite programs for each OS?

- POSIX API: a standard set of system calls that an OS must implement
- –Programs written to the POSIX API can run on any POSIX compliant OS
- –Most modern
- OSes are POSIX compliant–Ensures program portability
- Program language libraries hide the details of invoking system calls
- –The printf function in the C library calls the write system call to write to screen
- –User programs usually do not need to worry about invoking system calls

Summary

- Processes
 - At any given time, system has multiple active processes
 - On a one-CPU system, only one can execute at a time, but each process appears to have total control of the processor
 - OS periodically “context switches” between active processes
- Process management
 - `fork`: one call, two returns
 - `execve`: one call, usually no return
 - `wait` or `waitpid`: synchronization
 - `exit`: one call, no return

`fork` makes two copies of the same process (parent & child)

Returns different values to the two processes

`exec*` replaces current process from file (new program)

Two-process program:

- First `fork()`
- `if (pid == 0) { /* child code */ } else { /* parent code */ }`

Two different programs:

- First `fork()`
- `if (pid == 0) { execv(...) } else { /* parent code */ }`

`wait` or `waitpid` used to synchronize parent/child execution and to reap child process

Process Execution

- OS creates a process entry for the process in a process list
- Allocates memory and creates memory image
 - –Code and data (from executable)
 - –Stack and heap
- • Points CPU program counter to current instruction
 - –Other registers may store operands, return values etc.
- • After setup, OS is out of the way and process executes directly on CPU

A simple function call

- A function call translates to a jump instruction
- A new stack frame pushed to stack and stack pointer (SP) updated
- Old value of PC (return value) pushed to stack and PC updated
- Stack frame contains return value, function arguments etc.

How is a system call different?

- CPU hardware has multiple privilege levels
 - One to run user code: user mode
 - One to run OS code like system calls: kernel mode
 - Some instructions execute only in kernel mode
- Kernel does not trust user stack
 - Uses **a separate kernel stack** when in kernel mode
- Kernel does not trust user provided addresses to jump to
 - Kernel sets up **Interrupt Descriptor Table (IDT) at boot time**
 - IDT has addresses of kernel functions** to run for system calls and other events

Mechanism of system call: trap instruction

- When system call must be made, a special trap instruction is run(usually hidden from user by libc)
- Trap instruction execution
 - Move CPU to higher privilege level
 - Switch to kernel stack
 - Save context (old PC, registers) on kernel stack
 - Look up address in IDT and jump to trap handler function in OS code

More on trap instruction

- Trap instruction is executed on hardware in following cases:
 - System call (program needs OS service)
 - Program fault (program does something illegal, e.g., access memory it doesn't have access to)
 - Interrupt (external device needs attention of OS, e.g., a network packet has arrived on network card)
- Across all cases, the mechanism is: save context on kernel stack and switch to OS address in IDT
- IDT has many entries: which to use?
 - System calls/interrupts store a number in a CPU register before calling trap, to identify which IDT entry to use