

# Lecture 11

Jan 29, 2024

# Threads & Concurrency

---

So, far we have studied single threaded programs

- Recall: Process execution
  - PC points to current instruction being run
  - SP points to stack frame of current function call
- A program can also have multiple threads of execution
- What is a thread?

# Why do we need ?

---

Consider the case of a word processor :

It has several actions including a) respond to user's keystrokes and b) Spell check

Don't want to slow response to user's keystrokes because of spell checking

=> We need concurrency

These concurrent work-items must share the same data (text buffer), OS resources (eg, files, display functions etc)

# Why do we need ?

---


Consider the case of a web server:

Multiple requests need to be serviced concurrently


Especially important to maximize performance as one request may be held upon an I/O call

These request handlers share the same code, OS resources (eg. Sockets), and perhaps persistent stores like databases etc

# Multicores



Intel® Core™ i7-5960X



# Remember : What is an OS?

---

Middleware between user programs and system hardware that provides the following:

## Abstraction

OS makes HW  
easy-to-use and program  
by providing interfaces

## Virtualization

OS creates an illusion of  
dedicated HW for each user  
and application

## Concurrency

OS enables controlled  
interaction between  
multiple applications

# Do we know how to do that?

---

Till now, We know processes and IPCs

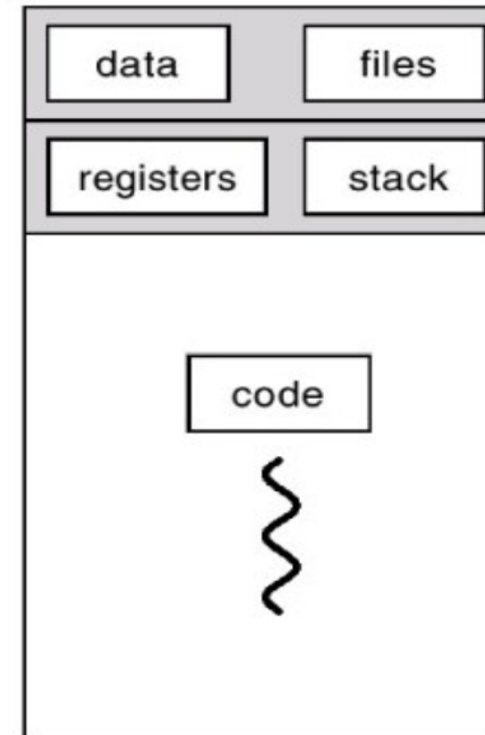
For the word processor application, we can create a new process that does spell check and communicates by files/signals

For the webserver, whenever we get a new request, we can fork a child that services the request

# Process Execution

---

- Separate streams of execution
- Each process isolated from the other
- Process state contains
  - Process ID
  - Environment
  - Working directory.
  - Program instructions
  - Registers
  - Stack
  - Heap
  - File descriptors
- Created by the OS using fork
  - Significant overheads





# Process is an overkill

---

In essence a process tracks three things

**Address Space** : Code, data, heap, (PT)

**Execution State** : PC, SP, register values ...

**OS State** : ID, open files, user permissions, directory, network connections, signals etc

Which of these do we need to share amongst the concurrent units from the earlier examples ?

# Process is an overkill

---

**Address Space** : Code, data, heap, (PT) -- SHARED

**Execution State** : PC, SP, register values -- SEPARATE

**OS State** : ID, open files, user permissions, directory, network connections, signals etc

**Need a 'lightweight process' as the unit of execution**

**"THREADS"**

# Threads


---

A thread is like another copy of a process that executes independently

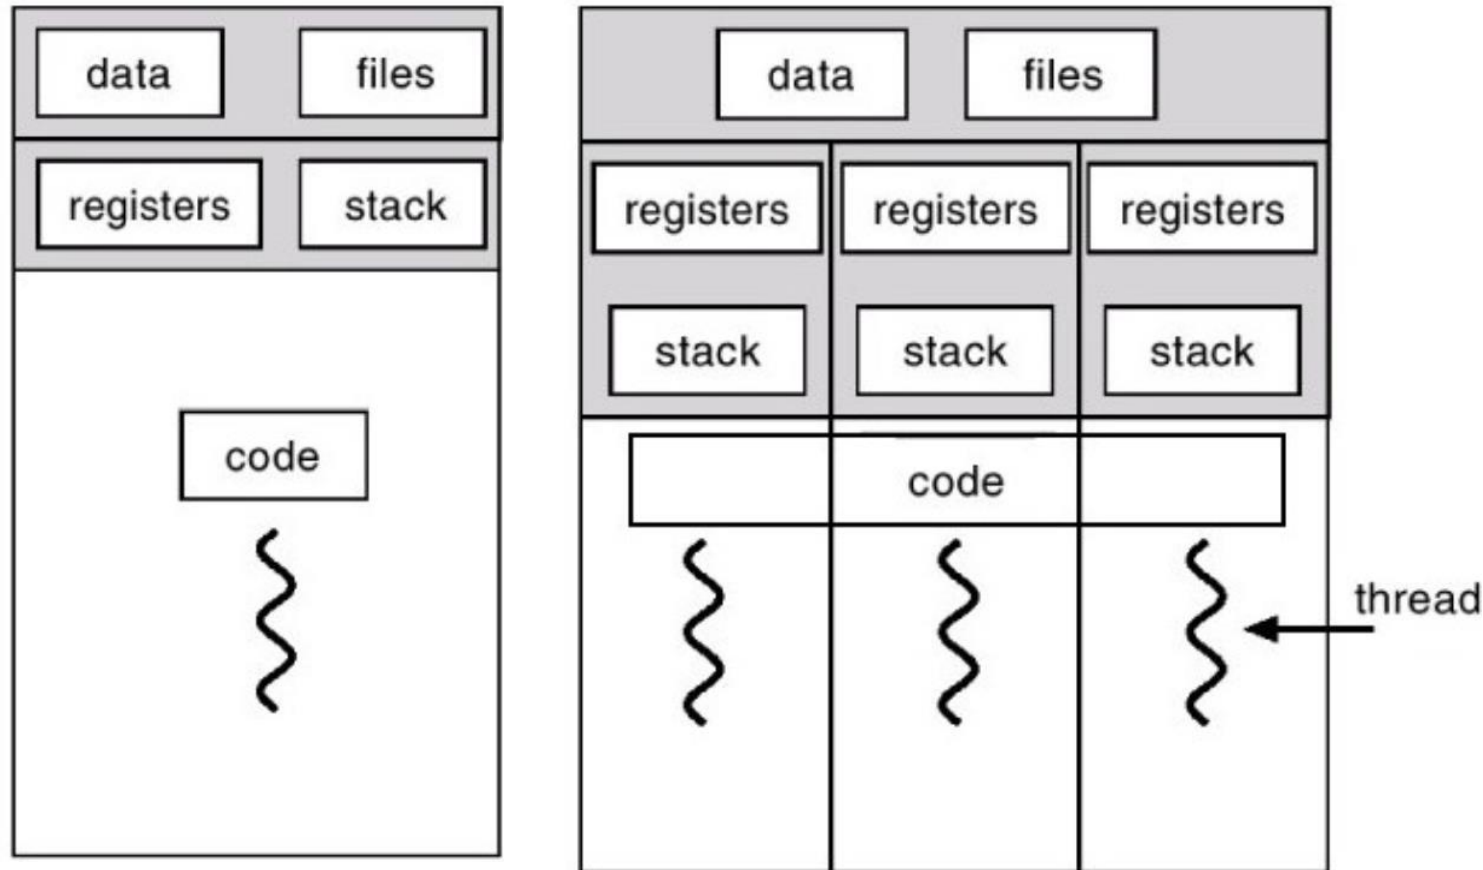
- Threads shares the same address space (code, heap)
- Each thread has separate PC
  - Each thread may run over different part of the program
- Each thread has separate stack for independent function calls

# Threads

- Separate streams of execution within a single process
- Threads in a process not isolated from each other
- Each thread state (thread control block) contains
  - Registers (including EIP, ESP)
  - stack



# Threads vs Processes



Key Idea: Separate **process info** (address space, OS resources) from **execution thread** (PC, regs, stack)

# Why threads ?

---

- Parallelism: a single process can effectively utilize multiple CPU cores
  - Understand the difference between concurrency and parallelism
  - Concurrency: running multiple threads/processes at the same time, even on single CPU core, by interleaving their executions
  - Parallelism: running multiple threads/processes in parallel over different CPU cores
- Even if no parallelism, concurrency of threads ensures effective use of CPU when one of the threads blocks (e.g., for I/O)

# Threads - Advantages

---

Communication between threads is very efficient – Same address space

Context switch between threads of same process is faster: Eg. No page table change

Each thread can work on its own data: Independent stack and registers

Threads can advantage of multi-core systems : Schedulable entity is a thread

# Why threads?

- Lightweight

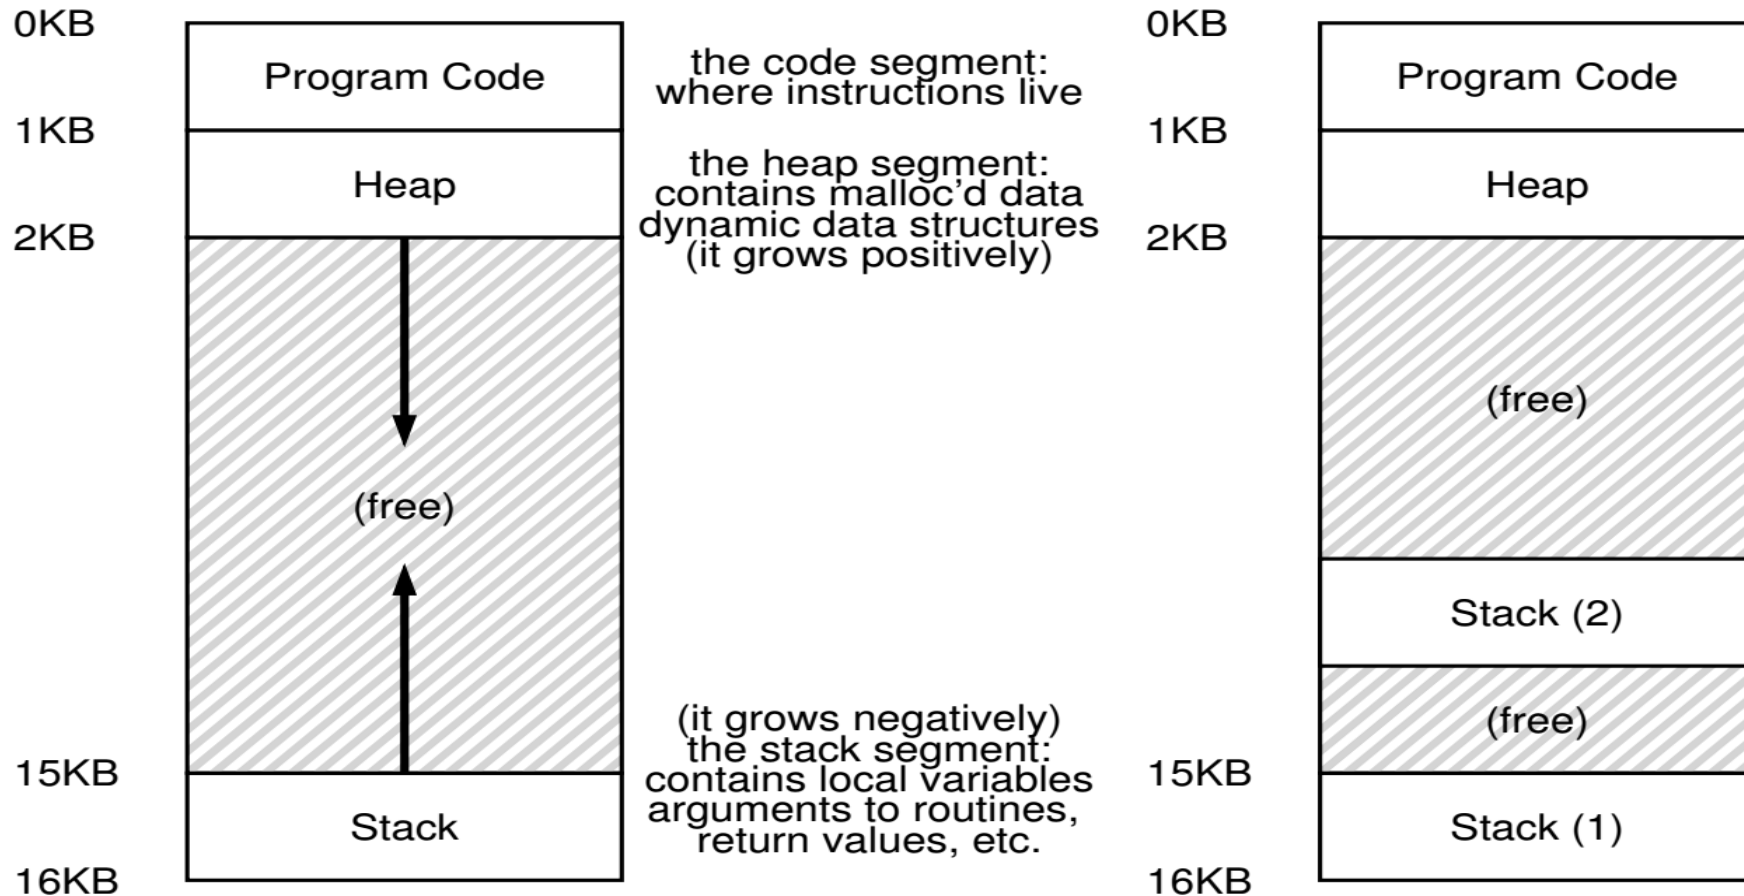
Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

Cost of creating 50,000 processes / threads  
(<https://computing.llnl.gov/tutorials/pthreads/>)

- Efficient communication between entities
- Efficient context switching



# Single-Threaded And Multi-Threaded Address Spaces



# Threads vs Processes

---

Parent P forks a child C

- P and C do not share any memory
- Need complicated IPC mechanisms to communicate
- Extra copies of code, data in memory
- Parent P executes two threads T1 and T2
  - T1 and T2 share parts of the address space
  - Global variables can be used for communication
  - Smaller memory footprint
- Threads are like separate processes, except they share the same address space

# Comparison

---

- A thread has no data segment or heap
- A thread cannot live on its own. It needs to be attached to a process
- There can be more than one thread in a process. Each thread has its own stack
- If a thread dies, its stack is reclaimed
- A process has code, heap, stack, other segments
- A process has at-least one thread.
- Threads within a process share the same I/O, code, files.
- If a process dies, all threads die.

# Trivia

---

Do the following statements apply to **Processes (P)**, **Threads (T)** or both (B)

Can share a virtual address space

Take longer to context switch

Have an execution context

Usually result in hotter caches when multiple exist

Make use of some communication mechanisms

# Scheduling Threads

---

- OS schedules threads that are ready to run independently, much like processes
- The context of a thread (PC, registers) is saved into/restored from thread control block (TCB)
  - Every PCB has one or more linked TCBs
- Threads that are scheduled independently by kernel are called kernel threads
  - E.g., Linux pthreads are kernel threads
- In contrast, some libraries provide user-level threads
  - User program sees multiple threads
  - Library multiplexes larger number of user threads over a smaller number of kernel threads
  - Low overhead of switching between user threads (no expensive context switch)
  - But multiple user threads cannot run in parallel

# pthread library

---

- Create a thread in a process

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine) (void *),  
                  void *arg);
```

Thread identifier (TID) much like

Pointer to a function,  
which starts execution in a  
different thread

Arguments to the function

- Destroying a thread

```
void pthread_exit(void *retval);
```

Exit value of the thread

# Pthread contd..

---

- Join : Wait for a specific thread to complete

```
int pthread_join(pthread_t thread, void **retval);
```

TID of the thread to wait for



Exit status of the thread



what is the difference with wait()?

```
#define Pthread_create(thread, attr, start_routine, arg) assert(pthread_create(thread, attr, start_routine, arg) == 0);
#define Pthread_join(thread, value_ptr) assert(pthread_join(thread, value_ptr) == 0);
```

# Example

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4  #include "common.h"
5  #include "common_threads.h"
6
7  void *mythread(void *arg) {
8      printf("%s\n", (char *) arg);
9      return NULL;
10 }
11
12 int
13 main(int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     Pthread_create(&p1, NULL, mythread, "A");
18     Pthread_create(&p2, NULL, mythread, "B");
19     // join waits for the threads to finish
20     Pthread_join(p1, NULL);
21     Pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }
```

```
#include <pthread.h>
#include <stdio.h>

void *thread_fn(void *arg){
    long id = (long) arg;
    printf("Starting thread %ld\n", id);
    sleep(5);
    printf("Exiting thread %ld\n", id);
    return NULL;
}

int main(){
    pthread_t t1, t2;

    pthread_create(&t1, NULL, thread_fn, (void *)1);
    pthread_create(&t2, NULL, thread_fn, (void *)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Exiting main\n");
    return 0;
}
```



# Thread Trace

main	Thread 1	Thread2	main	Thread 1	Thread2	main	Thread 1	Thread2
starts running			starts running			starts running		
prints "main: begin"			prints "main: begin"			prints "main: begin"		
creates Thread 1			creates Thread 1			creates Thread 1		
creates Thread 2				runs		creates Thread 2		
waits for T1				prints "A"				runs
	runs		creates Thread 2	returns				prints "B"
	prints "A"				runs	waits for T1		returns
	returns				prints "B"		runs	
waits for T2		runs	waits for T1		returns		prints "A"	
		prints "B"	<i>returns immediately; T1 is done</i>			waits for T2		returns
		returns	waits for T2			<i>returns immediately; T2 is done</i>		
prints "main: end"			prints "main: end"			prints "main: end"		

# Example: With Shared data

---

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include "common.h"
4  #include "common_threads.h"
5
6  static volatile int counter = 0;
7
8  // mythread()
9  //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *mythread(void *arg) {
15     printf("%s: begin\n", (char *) arg);
16     int i;
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
20     printf("%s: done\n", (char *) arg);
21     return NULL;
22 }
23
24 // main()
25 //
26 // Just launches two threads (pthread_create)
27 // and then waits for them (pthread_join)
28 //
29 int main(int argc, char *argv[]) {
30     pthread_t p1, p2;
31     printf("main: begin (counter = %d)\n", counter);
32     Pthread_create(&p1, NULL, mythread, "A");
33     Pthread_create(&p2, NULL, mythread, "B");
34
35     // join waits for the threads to finish
36     Pthread_join(p1, NULL);
37     Pthread_join(p2, NULL);
38     printf("main: done with both (counter = %d)\n",
39           counter);
40     return 0;
41 }
```

# With shared data – What happens?

---

```
prompt> gcc -o main main.c -Wall -pthread; ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 20000000)
```

What do we expect? Two threads, each increments counter by  $10^7$ , so  $2 \times 10^7$

```
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19345221)
```

Sometimes, a lower value. Why?

# Assembly code

OS	Thread 1	Thread 2	(after instruction)			Assembly code of
			PC	eax	counter	
	<i>before critical section</i>		100	0	50	counter = counter + 1
	mov 8049a1c, %eax		105	<b>50</b>	50	
	add \$0x1, %eax		108	<b>51</b>	50	100 mov 0x8049a1c, %eax
<b>interrupt</b>						105 add \$0x1, %eax
<i>save T1</i>						108 mov %eax, 0x8049a1c
<i>restore T2</i>			100	0	50	
		mov 8049a1c, %eax	105	<b>50</b>	50	
		add \$0x1, %eax	108	<b>51</b>	50	
		mov %eax, 8049a1c	113	51	<b>51</b>	
<b>interrupt</b>						
<i>save T2</i>						
<i>restore T1</i>			108	51	51	
	mov %eax, 8049a1c		113	51	<b>51</b>	

# Race condition & Synchronization

---

What just happened is called a race condition –Concurrent execution can lead to different results (functional output depends on order of execution)

Non-deterministic scheduling can change results

- Critical section:** portion of code that can lead to race conditions

Usually critical sections access shared resources (eg. Variable on heap)

- What we need: mutual exclusion–Only one thread should be executing critical section at any time

- What we need: atomicity of the critical section –**The critical section should execute like one uninterruptible instruction**