

# LECTURE 15

## FEB 7, 2024

# Condition Variables - Example

---

```
int done = 0;
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t c = PTHREAD_COND_INITIALIZER;

void thr_exit() {
    Pthread_mutex_lock(&m);
    done = 1;
    Pthread_cond_signal(&c);
    Pthread_mutex_unlock(&m);
}

void *child(void *arg) {
    printf("child\n");
    thr_exit();
    return NULL;
}

void thr_join() {
    Pthread_mutex_lock(&m);
    while (done == 0)
        Pthread_cond_wait(&c, &m);
    Pthread_mutex_unlock(&m);
}

int main(int argc, char *argv[]) {
    printf("parent: begin\n");
    pthread_t p;
    Pthread_create(&p, NULL, child, NULL);
    thr_join();
    printf("parent: end\n");
    return 0;
}
```

# Why while loop?

---

In the example code, why do we check condition before calling wait?—In case the child has already run and done is true, then no need to wait

- Why check condition with “while” loop and not “if”?
  - To avoid corner cases of thread being woken up even when condition not true (may be an issue with some implementations)

# Why use lock when calling wait?

---

What if no lock is held when calling wait/signal?

- Race condition: missed wakeup
  - Parent checks done to be 0, decides to sleep, gets interrupted
  - Child runs, sets done to 1, signals, but no one sleeping yet
  - Parent now resumes and goes to sleep forever
- Lock must be held when calling wait and signal with CV
- The wait function releases the lock before putting thread to sleep, so lock is available for signaling thread

```
void thr_exit() {  
    done = 1;  
    Pthread_cond_signal(&c);  
}
```

```
void thr_join() {  
    if (done == 0)  
        Pthread_cond_wait(&c);  
}
```

---


The responsibility of `wait()` is to release the lock and put the calling thread to sleep (atomically).

when the thread wakes up (after some other thread has signaled it), it must re-acquire the lock before returning to the caller.

# The Producer/Consumer Problem

---

- Also known as *Bounded buffer Problem*
- Producer produces and stores in buffer, Consumer consumes from buffer
- Trouble when
  - Producer produces, but buffer is full
  - Consumer consumes, but buffer is empty



# Producer/Consumer problem

---

- A common pattern in multi-threaded programs
- Example: in a multi-threaded web server, one thread accepts requests from the network and puts them in a queue. Worker threads get requests from this queue and process them.
- Setup: one or more producer threads, one or more consumer threads, a shared buffer of bounded size

# Attempt-1

---

```
int loops; // must initialize somewhere...
cond_t cond;
mutex_t mutex;
```

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex); // p1
        if (count == 1) // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i); // p4
        Pthread_cond_signal(&cond); // p5
        Pthread_mutex_unlock(&mutex); // p6
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex); // c1
        if (count == 0) // c2
            Pthread_cond_wait(&cond, &mutex); // c3
        int tmp = get(); // c4
        Pthread_cond_signal(&cond); // c5
        Pthread_mutex_unlock(&mutex); // c6
        printf("%d\n", tmp);
    }
}
```

```
int buffer;
int count = 0; // initially, empty
```

```
void put(int value) {
    assert(count == 0);
    count = 1;
    buffer = value;
}

int get() {
    assert(count == 1);
    count = 0;
    return buffer;
}
```



$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment	
c1	Run		Ready		Ready	0		int loops; // must initialize somewhere...
c2	Run		Ready		Ready	0		cond_t cond;
c3	Sleep		Ready		Ready	0	Nothing to get	mutex_t mutex;
	Sleep		Ready	p1	Run	0		void *producer(void *arg) {
	Sleep		Ready	p2	Run	0		int i;
	Sleep		Ready	p4	Run	1	Buffer now full	for (i = 0; i < loops; i++) {
	Ready		Ready	p5	Run	1	$T_{c1}$ awoken	pthread_mutex_lock(&mutex); // p1
	Ready		Ready	p6	Run	1		if (count == 1) // p2
	Ready		Ready	p1	Run	1		pthread_cond_wait(&cond, &mutex); // p3
	Ready		Ready	p2	Run	1		put(i); // p4
	Ready		Ready	p3	Sleep	1	Buffer full; sleep	pthread_cond_signal(&cond); // p5
	Ready	c1	Run		Sleep	1	$T_{c2}$ sneaks in ...	pthread_mutex_unlock(&mutex); // p6
	Ready	c2	Run		Sleep	1		}
	Ready	c4	Run		Sleep	0	... and grabs data	void *consumer(void *arg) {
	Ready	c5	Run		Ready	0	$T_p$ awoken	int i;
	Ready	c6	Run		Ready	0		for (i = 0; i < loops; i++) {
c4	Run		Ready		Ready	0	Oh oh! No data	pthread_mutex_lock(&mutex); // c1
								if (count == 0) // c2
								pthread_cond_wait(&cond, &mutex); // c3
								int tmp = get(); // c4
								pthread_cond_signal(&cond); // c5
								pthread_mutex_unlock(&mutex); // c6
								printf("%d\n", tmp);
								}

# Attempt-2

---

```
int loops;
cond_t  cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // p1
        while (count == 1)                    // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i);                               // p4
        Pthread_cond_signal(&cond);           // p5
        Pthread_mutex_unlock(&mutex);         // p6
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // c1
        while (count == 0)                    // c2
            Pthread_cond_wait(&cond, &mutex); // c3
        int tmp = get();                      // c4
        Pthread_cond_signal(&cond);           // c5
        Pthread_mutex_unlock(&mutex);         // c6
        printf("%d\n", tmp);
    }
}
```

# Thread Trace

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment	
c1	Run		Ready		Ready	0		int loops;
c2	Run		Ready		Ready	0		cond_t cond;
c3	Sleep		Ready		Ready	0	Nothing to get	mutex_t mutex;
	Sleep	c1	Run		Ready	0		void *producer(void *arg) {
	Sleep	c2	Run		Ready	0		int i;
	Sleep	c3	Sleep		Ready	0	Nothing to get	for (i = 0; i < loops; i++) {
	Sleep		Sleep	p1	Run	0		pthread_mutex_lock(&mutex); // p1
	Sleep		Sleep	p2	Run	0		while (count == 1) // p2
	Sleep		Sleep	p4	Run	1	Buffer now full	pthread_cond_wait(&cond, &mutex); // p3
	Ready		Sleep	p5	Run	1	$T_{c1}$ awoken	put(i); // p4
	Ready		Sleep	p6	Run	1		pthread_cond_signal(&cond); // p5
	Ready		Sleep	p1	Run	1		pthread_mutex_unlock(&mutex); // p6
	Ready		Sleep	p2	Run	1		}
	Ready		Sleep	p3	Sleep	1	Must sleep (full)	void *consumer(void *arg) {
c2	Run		Sleep		Sleep	1	Recheck condition	int i;
c4	Run		Sleep		Sleep	0	$T_{c1}$ grabs data	for (i = 0; i < loops; i++) {
c5	Run		Ready		Sleep	0	Oops! Woke $T_{c2}$	pthread_mutex_lock(&mutex); // c1
c6	Run		Ready		Sleep	0		while (count == 0) // c2
c1	Run		Ready		Sleep	0		pthread_cond_wait(&cond, &mutex); // c3
c2	Run		Ready		Sleep	0		int tmp = get(); // c4
c3	Sleep		Ready		Sleep	0	Nothing to get	pthread_cond_signal(&cond); // c5
	Sleep	c2	Run		Sleep	0		pthread_mutex_unlock(&mutex); // c6
	Sleep	c3	Sleep		Sleep	0	Everyone asleep...	printf("%d\n", tmp);
								}
								}

# Attempt-3

---

```
cond_t empty, fill;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 1)
            Pthread_cond_wait(&empty, &mutex);
        put(i);
        Pthread_cond_signal(&fill);
        Pthread_mutex_unlock(&mutex);
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 0)
            Pthread_cond_wait(&fill, &mutex);
        int tmp = get();
        Pthread_cond_signal(&empty);
        Pthread_mutex_unlock(&mutex);
        printf("%d\n", tmp);
    }
}
```

# Solution

---

```
int buffer[MAX];
int fill_ptr = 0;
int use_ptr = 0;
int count = 0;

void put(int value) {
    buffer[fill_ptr] = value;
    fill_ptr = (fill_ptr + 1) % MAX;
    count++;
}

int get() {
    int tmp = buffer[use_ptr];
    use_ptr = (use_ptr + 1) % MAX;
    count--;
    return tmp;
}
```

```
cond_t empty, fill;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // p1
        while (count == MAX)                 // p2
            Pthread_cond_wait(&empty, &mutex); // p3
        put(i);                             // p4
        Pthread_cond_signal(&fill);          // p5
        Pthread_mutex_unlock(&mutex);        // p6
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // c1
        while (count == 0)                   // c2
            Pthread_cond_wait(&fill, &mutex); // c3
        int tmp = get();                    // c4
        Pthread_cond_signal(&empty);         // c5
        Pthread_mutex_unlock(&mutex);        // c6
        printf("%d\n", tmp);
    }
}
```

# Semaphore

---

- Synchronization primitive like condition variables
- Semaphore is a variable with an underlying counter
- Two functions on a semaphore variable

```
#include <semaphore.h>
sem_t s;
sem_init(&s, 0, 1);
```

–**Up/post** increments the counter and wakes up one of the processes sleeping/blocked on the semaphore

–**Down/wait** decrements the counter and blocks the calling thread if the resulting value is negative

- 
- A semaphore with init value 1 acts as a simple lock(binary semaphore =mutex)

```
sem_t m;  
sem_init(&m, 0, X); // initialize to X; what should X be?  
  
sem_wait(&m);  
// critical section here  
sem_post(&m);
```

# POSIX semaphores

---

- `sem_init`
- `sem_wait`
- `sem_post`
- `sem_getvalue`
- `sem_destroy`



# Semaphores for ordering

---

- Can be used to set order of execution between threads like CV

- Example: parent waiting for child (init = 0)

```
sem_t s;

void *child(void *arg) {
    printf("child\n");
    sem_post(&s); // signal here: child is done
    return NULL;
}

int main(int argc, char *argv[]) {
    sem_init(&s, 0, X); // what should X be?
    printf("parent: begin\n");
    pthread_t c;
    Pthread_create(&c, NULL, child, NULL);
    sem_wait(&s); // wait here for child
    printf("parent: end\n");
    return 0;
}
```

# Producer – Consumer Problem

---

- Need two semaphores for signaling
  - One to track empty slots, and make producer wait if no more empty slots
  - One to track full slots, and make consumer wait if no more full slots
- One semaphore to act as mutex for buffer


---

Buffer of size N  
int count;

**full = 0, empty = N**

Number of filled blocks in  
the buffer

Number of empty blocks in  
the buffer



Buffer (of size N)

N = 6  
fill = 4  
empty = 2

full = 0, empty = N

```
void producer(){  
  while(TRUE){  
    → item = produce_item();  
    → down(empty);  
    → insert_item(item); // into buffer  
    → up(full);  
  }  
}
```



Buffer (of size N)

N = 6  
fill = 5  
empty = 1



Buffer (of size N)

N = 6

fill = 5

empty = 1

```
void consumer(){
```

```
  while(TRUE){
```

```
    → down(full);
```

```
    → item = remove_item(); // from buffer
```

```
    → up(empty);
```

```
    → consume_item(item);
```

```
  }
```

```
}
```

## The FULL Buffer

```
void producer(){
    while(TRUE){
        item = produce_item();
        → down(empty);
        → insert_item(item); // into buffer
        → up(full);
    }
}
```

```
void consumer(){
    while(TRUE){
        down(full);

        item = remove_item(); // from buffer

        up(empty);
        consume_item(item);
    }
}
```



Buffer (of size N)

N = 6  
fill = 6

**empty = 0**

## The Empty Buffer

```
void producer(){
    while(TRUE){
        item = produce_item();
        down(empty);

        insert_item(item); // into buffer

        up(full);
    }
}
```

```
void consumer(){
    while(TRUE){
        down(full);

        item = remove_item(); // from buffer

        up(empty);
        consume_item(item);
    }
}
```



Buffer (of size N)

N = 6

fill = 0

empty = 6

## Serializing Access to the Buffer

```
void producer(){
    while(TRUE){
        item = produce_item();
        down(empty);
        lock(mutex)
        insert_item(item); // into buffer
        unlock(mutex)
        up(full);
    }
}
```

```
void consumer(){
    while(TRUE){
        down(full);
        lock(mutex)
        item = remove_item(); // from buffer
        unlock(mutex)
        up(empty);
        consume_item(item);
    }
}
```



Buffer (of size N)

N = 6

fill = 3

empty = 3



---

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty);          // Line P1
        sem_wait(&mutex);          // Line P1.5 (MUTEX HERE)
        put(i);                    // Line P2
        sem_post(&mutex);          // Line P2.5 (AND HERE)
        sem_post(&full);           // Line P3
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&full);           // Line C1
        sem_wait(&mutex);          // Line C1.5 (MUTEX HERE)
        int tmp = get();           // Line C2
        sem_post(&mutex);          // Line C2.5 (AND HERE)
        sem_post(&empty);          // Line C3
        printf("%d\n", tmp);
    }
}

int main(int argc, char *argv[]) {
    // ...
    sem_init(&empty, 0, MAX); // MAX are empty
    sem_init(&full, 0, 0);    // 0 are full
    // ...
}
```

# Deadlock?

---

What if lock is acquired before signaling?

- Waiting thread sleeps with mutex and the signaling thread can never wake it up

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);           // Line P0 (NEW LINE)
        sem_wait(&empty);           // Line P1
        put(i);                     // Line P2
        sem_post(&full);             // Line P3
        sem_post(&mutex);           // Line P4 (NEW LINE)
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);           // Line C0 (NEW LINE)
        sem_wait(&full);            // Line C1
        int tmp = get();            // Line C2
        sem_post(&empty);           // Line C3
        sem_post(&mutex);           // Line C4 (NEW LINE)
        printf("%d\n", tmp);
    }
}
```